

Skill-it! A data-driven skills framework for understanding and training language models

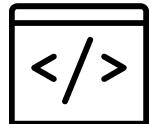
Mayee F. Chen, Nicholas Roberts, Kush Bhatia, Jue Wang, Ce Zhang, Fred Sala,
Christopher Ré

August 31st, Allen AI
mfchen@stanford.edu



Motivation

Large language models (LLMs) can do many things:



Write code



Chat with users

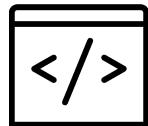


Generate creative content

Why? Architecture, training, **data**. How to select data?

Motivation

Large language models (LLMs) can do many things:



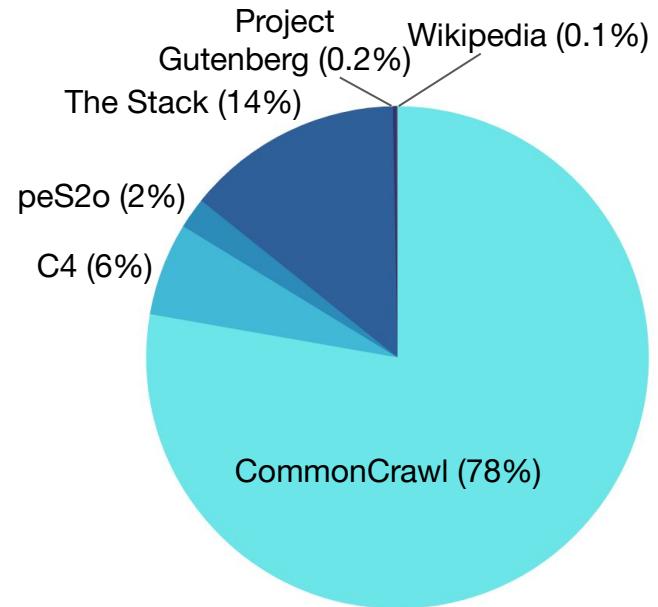
Write code



Chat with users



Generate creative content

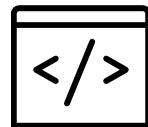


Why? Architecture, training, **data**. How to select data?

Training data mixture¹

Motivation

Large language models (LLMs) can do many things:



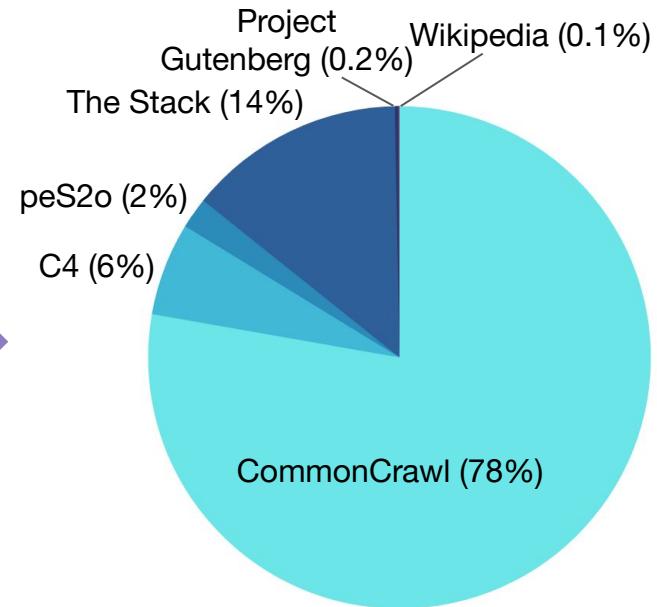
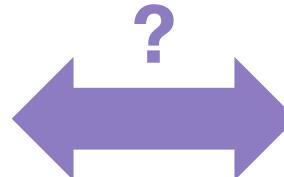
Write code



Chat with users



Generate creative content



Why? Architecture, training, **data**. How to select data?

Training data mixture¹

How is training data related to LLM's capabilities?

How do humans learn from data? Learn skills in a certain *order*.

How is training data related to LLM's capabilities?

How do humans learn from data? Learn skills in a certain *order*.



How is training data related to LLM's capabilities?

How do humans learn from data? Learn skills in a certain *order*.

Learning hierarchy:

- Learning addition also helps with: subtraction, multiplication, linear equations
- Learning subtraction also helps with: division, linear equations
- And so on

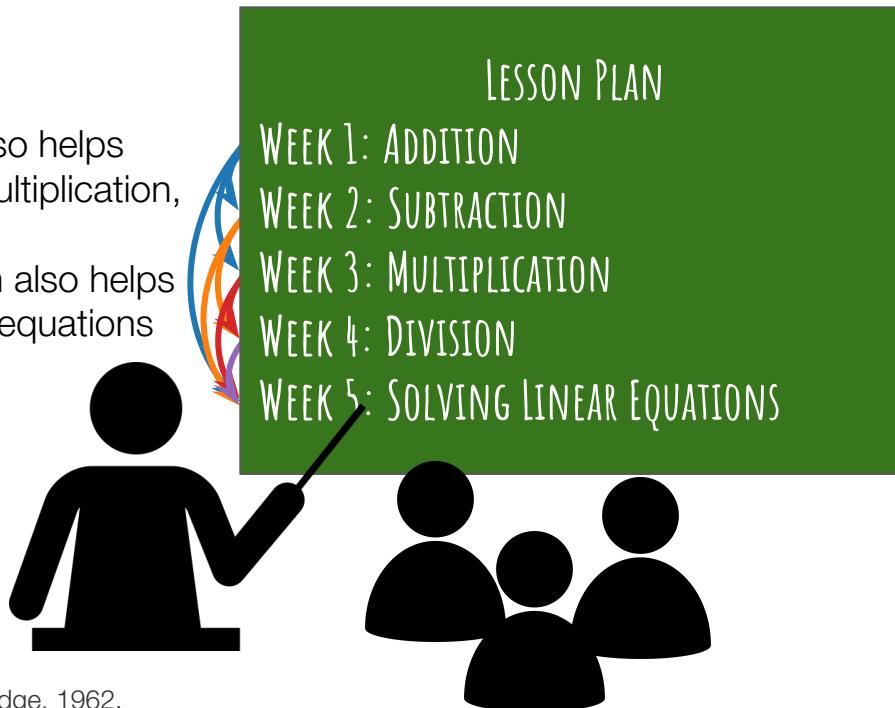


How is training data related to LLM's capabilities?

How do humans learn from data? Learn skills in a certain *order*.

Learning hierarchy:

- Learning addition also helps with: subtraction, multiplication, linear equations
- Learning subtraction also helps with: division, linear equations
- And so on



Changing the order of skills makes learning harder!¹

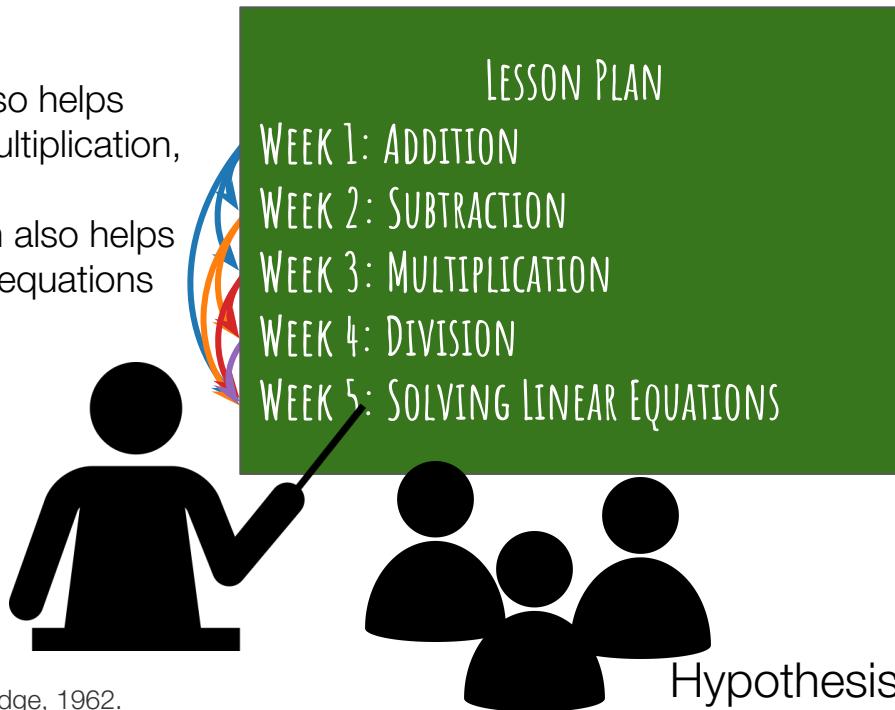
- Learn division then addition → students do worse

[1] Gagné. The acquisition of knowledge, 1962.

How is training data related to LLM's capabilities?

How do humans learn from data? Learn skills in a certain *order*.

- Learning addition also helps with: subtraction, multiplication, linear equations
- Learning subtraction also helps with: division, linear equations
- And so on



Changing the order of skills makes learning harder!¹

- Learn division then addition
→ students do worse

Hypothesis: models also learn like this

[1] Gagné. The acquisition of knowledge, 1962.

Q: How does training data influence various model capabilities?

Can we use this understanding to more effectively select data to improve such capabilities?

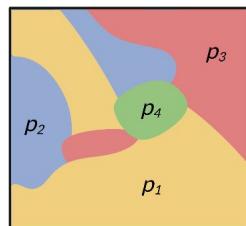
Q: How does training data influence various model capabilities?

Can we use this understanding to more effectively select data to improve such capabilities?

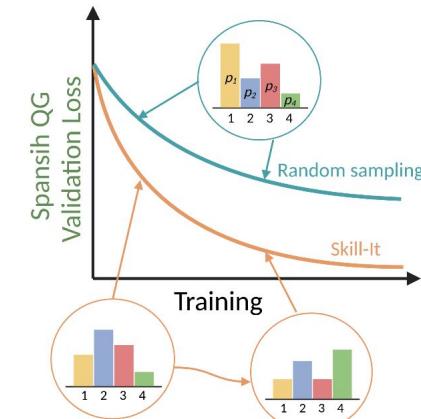
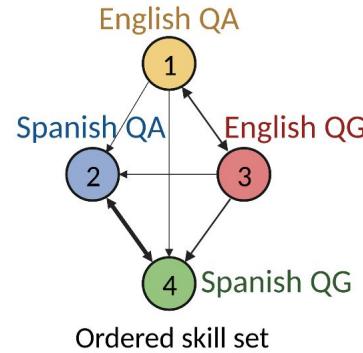
A: there exist sets of *skills* associated with data that the LLM learns most efficiently in some particular order. We can learn this order and exploit it to better select training data.

Outline

- Defining ordered skill sets
- Skill-It data selection algorithm
- Results
- Discussion



Data



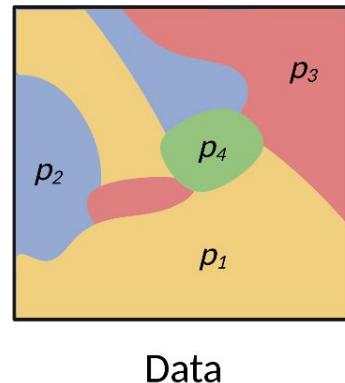
Defining ordered skill sets

Definitions: what is a skill?

Training data can be partitioned into subsets associated with *skills*.

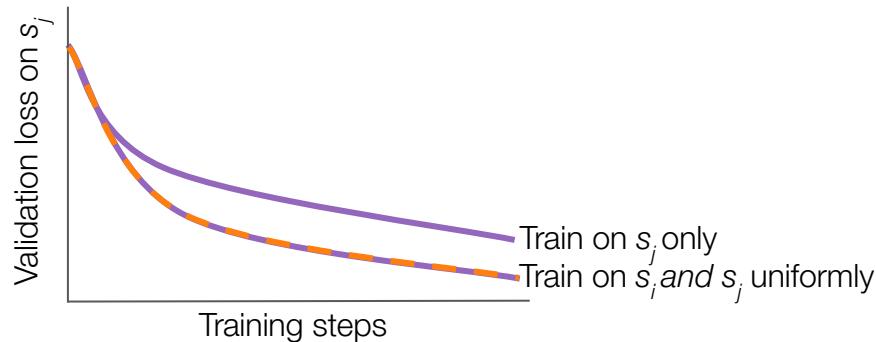
Def (informal): a **skill** s is a unit of behavior with associated data X_s s.t. if a model f is trained on dataset $D_s \subset X_s$, f has improved metric L (e.g., validation loss) on samples belonging to $X_s \setminus D_s$ on average.

Examples: tasks, data sources, task categories.



How do we define a meaningful order over skills?

Def: Given a set of skills S , its **skills graph** is $G = (S, E)$, where given a fixed data/training budget, $s_i \rightarrow s_j \in E$ iff validation loss on s_j when trained on mixture of s_i and s_j is no greater than when trained on just s_j .

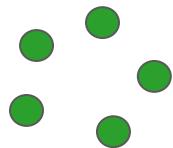


How do we define a meaningful order over skills?

Def: an **ordered skills set** is a set of skills whose **skills graph** is neither empty nor complete.

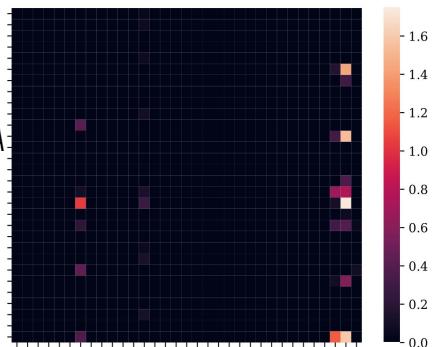
How do we define a meaningful order over skills?

Def: an **ordered skills set** is a set of skills whose **skills graph** is neither empty nor complete.



Dataset/skills: Pile of Law¹/data sources

Skills graph's
adjacency matrix A

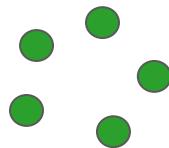


Ordered skill set?

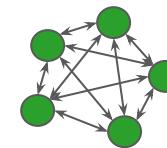


How do we define a meaningful order over skills?

Def: an **ordered skills set** is a set of skills whose **skills graph** is neither empty nor complete.

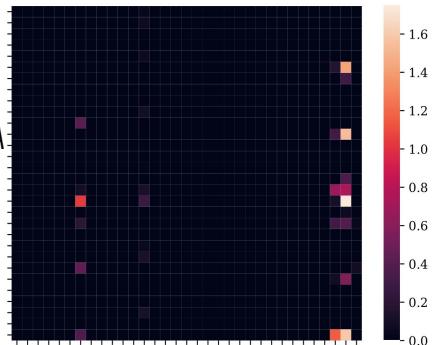


Dataset/skills: Pile of Law¹/data sources

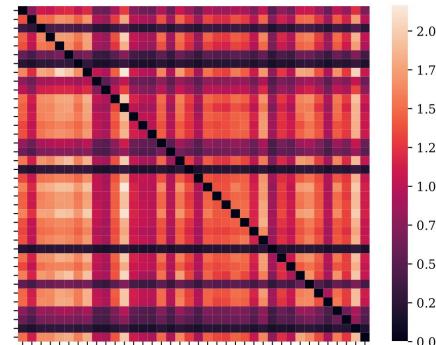


Alpaca³/leading verb

Skills graph's
adjacency matrix A

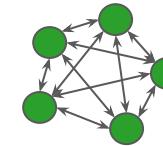
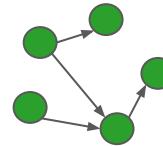
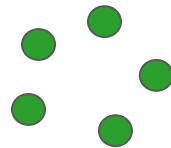


Ordered skill set?



How do we define a meaningful order over skills?

Def: an **ordered skills set** is a set of skills whose **skills graph** is neither empty nor complete.



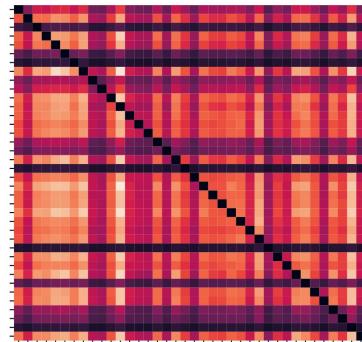
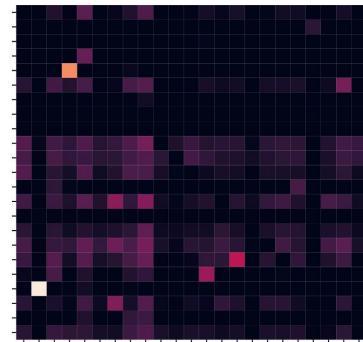
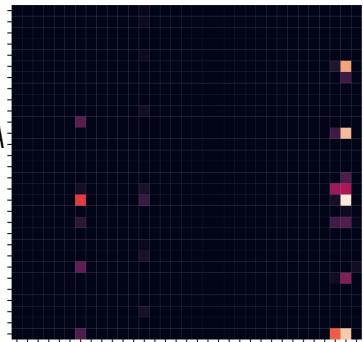
Dataset/skills:

Pile of Law¹/data sources

Natural Instructions²/task categories

Alpaca³/leading verb

Skills graph's
adjacency matrix A

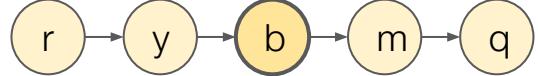
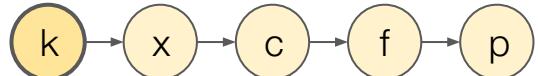


[1] Henderson et. al., 2022.

[2] Wang et. al., 2022.

[3] Taori et. al., 2023.

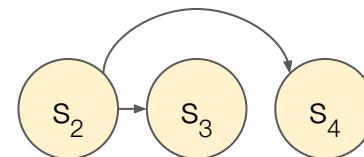
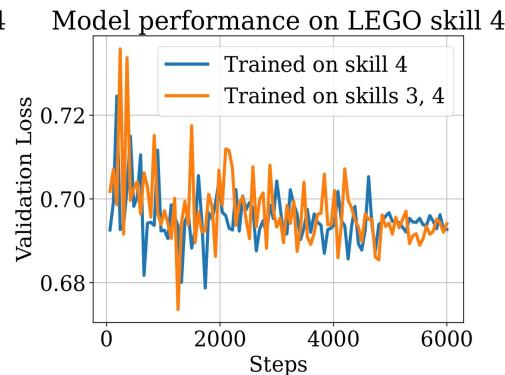
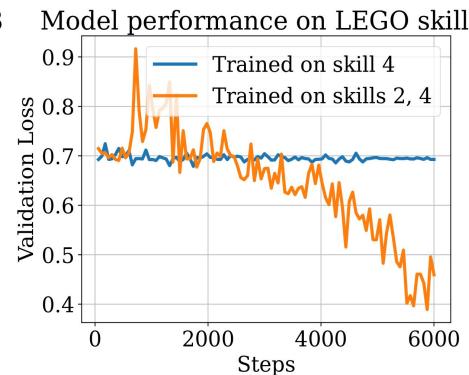
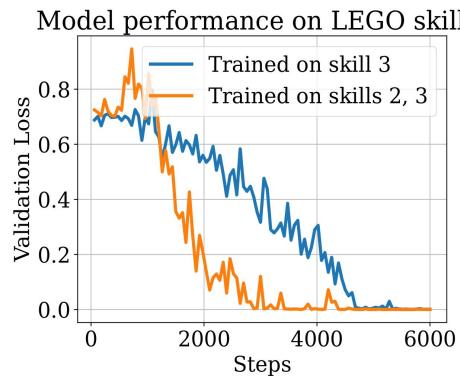
Ordered skill set examples: LEGO synthetic¹

Input	Output	
$b = \text{not } y, r = \text{val } 1, m = \text{val } b, q = \text{val } m, y = \text{not } r. b = ?$	1	 Skill 3
$c = \text{val } x, p = \text{val } f, x = \text{val } k, f = \text{not } c, k = \text{val } 0. k = ?$	0	 Skill 1

- Each sample is a series of recursive variable assignments to 0 or 1. The model is asked to provide the value of a variable.
- Skill X = model's ability to get the value of X th variable correct
 - Skill 1 = recall only, easy
 - Skill k = need to learn skill $k-1$, harder
- Intuitive guess: skills graph is a “chain”

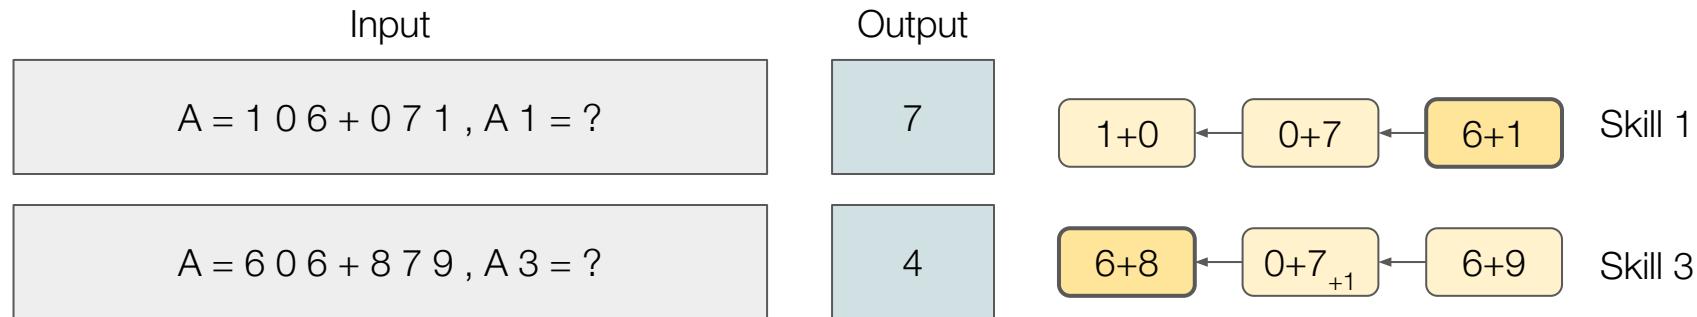
Ordered skill set examples: LEGO synthetic

- Continually pre-train GPT-Neo 125M on concatenated LEGO input/output pairs
- Measure loss on output token for held-out validation set



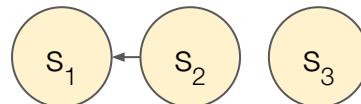
There is an ordered
skill set, but not the
one we guessed!

Ordered skill set examples: Addition synthetic



- Skill X = model's ability to get the Xth digit of the addition correct
 - Skill 1 = no carry, easy
 - Skill k = may need carry/dependent on skill $k-1$, harder
- Intuitive guess: skills graph is a “chain”

Ordered skill set examples: Addition synthetic



There is an ordered skill set, but not the one we guessed!

Ordered skill set examples: Natural Instructions

Skills

English question answering

Input: 3sat began broadcasting on 1 December 1984. What year did 3sat start?
Output: 1 december 1984

English question generation

Input: Snow Wolf is an espionage novel by Glenn Meade.
Output: Who wrote the Snow Wolf

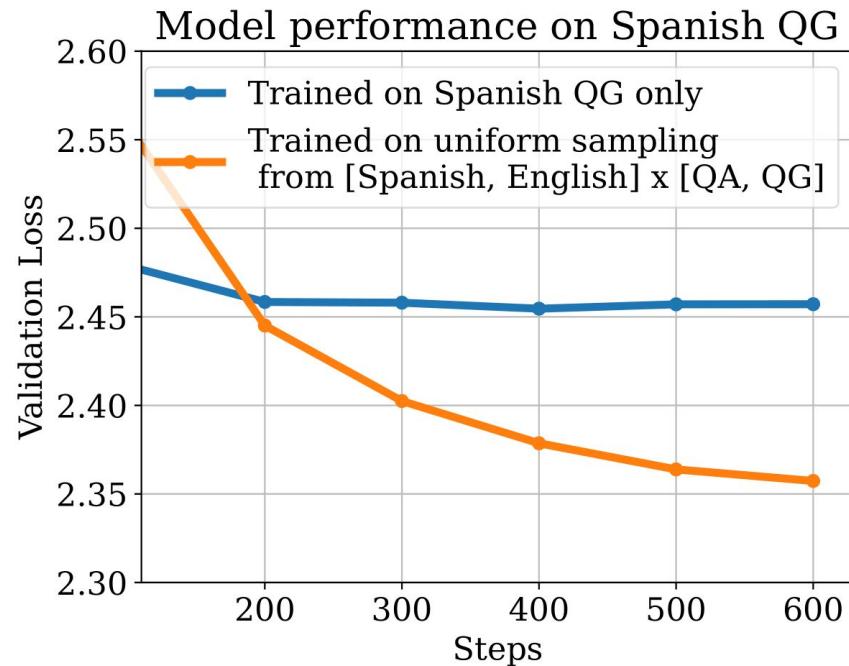
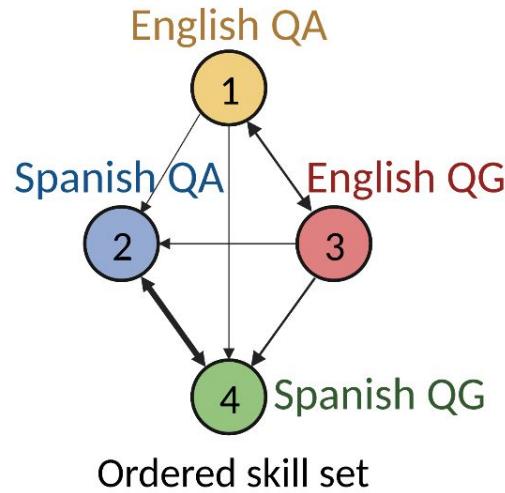
Spanish question answering

Input: Que inventen ellas es el título de la muestra sobre científicas e investigadoras, que estará hasta el 4 de enero en la Obra Social de Caja Madrid, en la Plaza de Catalunya. ¿Dónde estará disponible la muestra?
Output: en la Obra Social de Caja Madrid

Spanish question generation

Input: O sea, lloraba el músico. Volverá en octubre Pedrito. Volverá, ya lo verán. Volverá. Dios mío.
Output: ¿Cuándo volverá Pedrito?

Ordered skill set examples: real data



Ordered skill set examples: real data

Skills

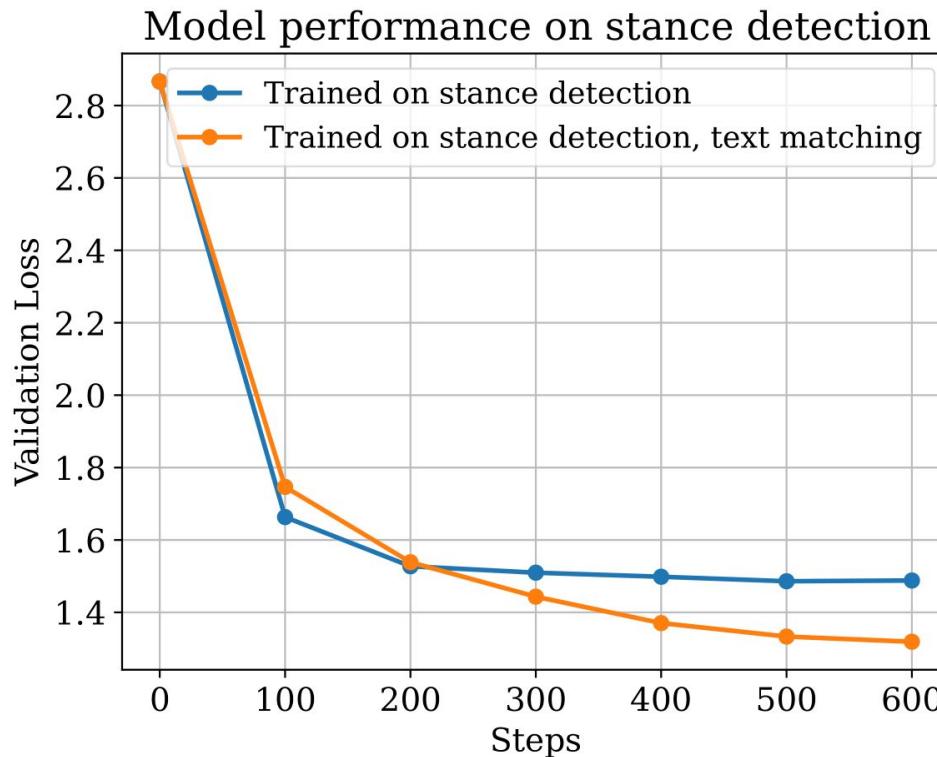
Stance Detection

Input:
Topic: Airport security profiling
Argument: Profiling will help avoid invasive scanners and pat-downs.
Output: in favor

Text Matching

Input:
Sentence 1: What are considered the safest prescription pills for high blood pressure/hypertension?
Sentence 2: What are the common side effects of blood pressure medications?
Output: Dissimilar

Ordered skill set examples: real data



Skill-It data selection algorithm

Problem Setup

Input:

- Training skill set $S_{\text{train}} = s_{\text{train},1}, \dots, s_{\text{train},k}$ with associated training corpus $X_{\text{train}} = X_{\text{train},1}, \dots, X_{\text{train},k}$
- Evaluation skill set $S_{\text{eval}} = s_{\text{eval},1}, \dots, s_{\text{eval},m}$ with associated held-out validation data X_{eval}
- Budget of n samples
- Pre-trained LLM f

Goal: how to order/select n samples from X_{train} for f to perform well on X_{eval}

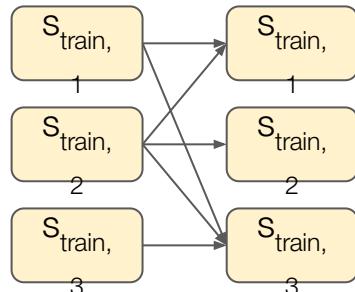
Problem Setting

We study three regimes depending on the relationship between S_{train} and S_{eval} :

Problem Setting

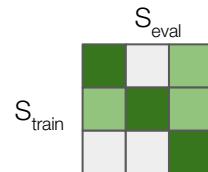
We study three regimes depending on the relationship between S_{train} and S_{eval} :

$$S_{\text{train}} = S_{\text{eval}}$$



Continual pre-training

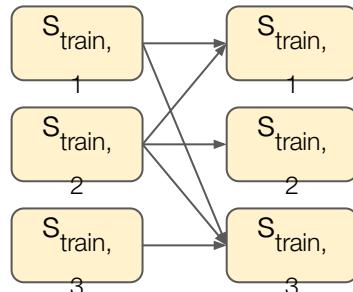
Adjacency
matrix A of
skills graph



Problem Setting

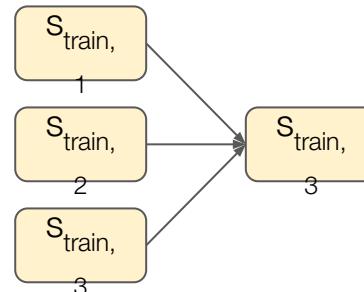
We study three regimes depending on the relationship between S_{train} and S_{eval} :

$$S_{\text{train}} = S_{\text{eval}}$$



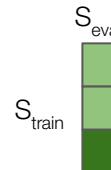
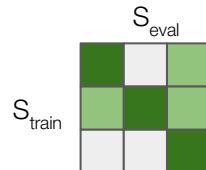
Continual pre-training

$$S_{\text{eval}} \subset S_{\text{train}}$$



Fine-tuning

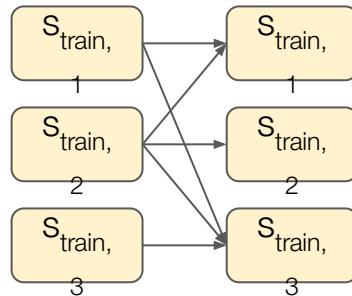
Adjacency matrix A of skills graph



Problem Setting

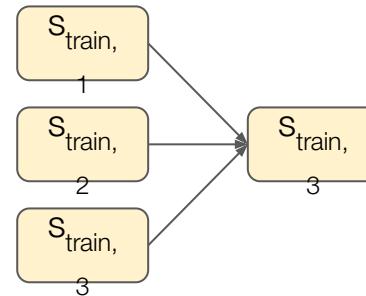
We study three regimes depending on the relationship between S_{train} and S_{eval} :

$$S_{\text{train}} = S_{\text{eval}}$$



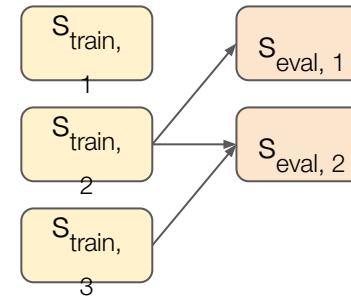
Continual pre-training

$$S_{\text{eval}} \subset S_{\text{train}}$$



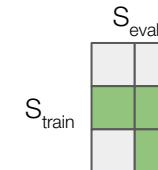
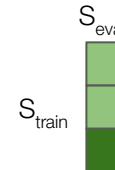
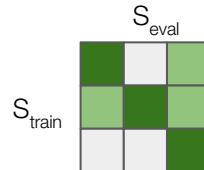
Fine-tuning

$$S_{\text{train}} \cap S_{\text{eval}} = \emptyset$$



Out-of-domain

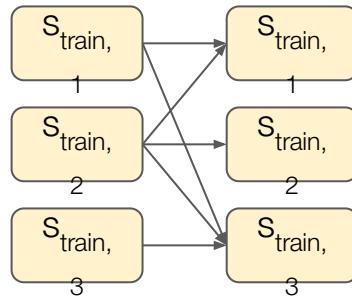
Adjacency matrix A of
skills graph



Problem Setting

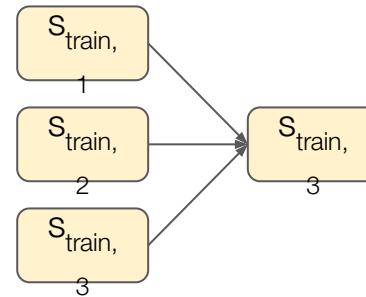
We study three regimes depending on the relationship between S_{train} and S_{eval} :

$$S_{\text{train}} = S_{\text{eval}}$$



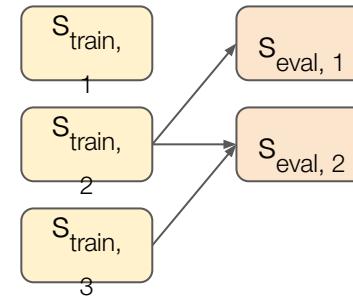
Continual pre-training

$$S_{\text{eval}} \subset S_{\text{train}}$$



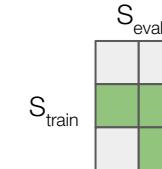
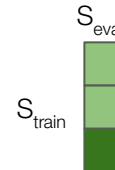
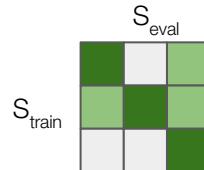
Fine-tuning

$$S_{\text{train}} \cap S_{\text{eval}} = \emptyset$$



Out-of-domain

Adjacency matrix A of skills graph



- Our general approach:
1. Learn A
 2. Sample according to A

Skills Graph Learning

Can follow the definition:

- For each s_i, s_j , see if validation loss on s_j is lower when trained on s_i and s_j than when trained on just s_j for the same amount of steps, H
 - Set edge weight proportional to change in loss
 - Runtime: $O(Hkm)$

Skills Graph Learning

Can follow the definition:

- For each s_i, s_j , see if validation loss on s_j is lower when trained on s_i and s_j than when trained on just s_j for the same amount of steps, H
 - Set edge weight proportional to change in loss
 - Runtime: $O(Hkm)$

We can also make this cheaper!

- Linear approximation: for each s_i, s_j , see if validation loss on s_j decreases when trained on s_i
 - Set edge weight proportional to change in loss
 - Runtime: $O(Hk)$
- Reduce number of steps H
- Learn skills graph using smaller model

Sampling according to Skills Graph

Def: **skill-stratified sampling** involves sampling uniform from all skills that are prerequisite for the evaluation skill set.

- $S' = \{s_{\text{train},i} \text{ if } \exists j \text{ s.t. } (s_{\text{train},i}, s_{\text{eval},j}) \in E\}$
- $\Pr(s_{\text{train},i}) = 1/|S'|$ if $s_{\text{train},i} \in S'$

Drawbacks: does not exploit skills graph dynamically - even after a skill is learned, we do not adjust skill proportions

Skill-It optimization problem

Formulate data selection as online optimization problem:

- T rounds: at each round t we sample n/T samples from X_{train} according to $p_t \in \Delta^{k-1}$ over the training skill set

Skill-It optimization problem

Formulate data selection as online optimization problem:

- T rounds: at each round t we sample n/T samples from X_{train} according to $p_t \in \Delta^{k-1}$ over the training skill set
- Let f_t be model at round t , with learning dynamics expressed as $f_t = \Phi(f_{t-1}, p_{t-1})$

Skill-It optimization problem

Formulate data selection as online optimization problem:

- T rounds: at each round t we sample n/T samples from X_{train} according to $p_t \in \Delta^{k-1}$ over the training skill set
- Let f_t be model at round t , with learning dynamics expressed as $f_t = \Phi(f_{t-1}, p_{t-1})$
- Let $L_{eval, j}(f_t)$ be the validation loss of f_t on $s_{eval, j}$.

Skill-It optimization problem

Formulate data selection as online optimization problem:

- T rounds: at each round t we sample n/T samples from X_{train} according to $p_t \in \Delta^{k-1}$ over the training skill set
- Let f_t be model at round t , with learning dynamics expressed as $f_t = \Phi(f_{t-1}, p_{t-1})$
- Let $L_{eval,j}(f_t)$ be the validation loss of f_t on $s_{eval,j}$.
- Objective: minimize average validation loss on evaluation skills at time T :

$$\begin{aligned} & \underset{p_1, \dots, p_T \in \Delta^{k-1}}{\text{minimize}} \quad \frac{1}{m} \sum_{j=1}^m L_{eval,j}(f_T) \\ & \text{s.t. } f_t = \Phi(f_{t-1}, p_{t-1}) \quad \forall t \in 1, \dots, T \end{aligned}$$

Skill-It optimization problem

- Hard to solve without more constraints/information on Φ
- We think about learning dynamics in terms of how loss of model f_t depends on loss of f_{t-1} and proportions p_{t-1}

Skill-It optimization problem

- Hard to solve without more constraints/information on Φ
- We think about learning dynamics in terms of how loss of model f_t depends on loss of f_{t-1} and proportions p_{t-1}
- Simple idea (assume $S_{\text{train}} = S_{\text{eval}}$ for now):

$$L_{\text{eval},j}(f_t) = L_{\text{eval},j}(\Phi(f_{t-1}, p_{t-1})) := L_{\text{eval},j}(f_{t-1})(1 - \alpha p_{t-1}^j)$$

- Loss on j at time t is loss on j at time $t-1$ decreased by a factor proportional to p_{t-1}^j .
 - More training data from $X_{\text{train},j}$ = loss goes down more
- Does not account for skills graph

Skill-It optimization problem

$$L_{\text{eval},j}(f_t) = L_{\text{eval},j}(\Phi(f_{t-1}, p_{t-1})) := L_{\text{eval},j}(f_{t-1})(1 - \alpha p_{t-1}^j)$$

Skill-It optimization problem

$$L_{\text{eval},j}(f_t) = L_{\text{eval},j}(\Phi(f_{t-1}, p_{t-1})) := L_{\text{eval},j}(f_{t-1})(1 - \boxed{A_{:,j}^\top p_{t-1}})$$

- More training data from $X_{\text{train}, j}$ or **any of j 's prerequisite skills** = loss on j goes down more

Skill-It optimization problem

$$\begin{aligned} & \underset{p_1, \dots, p_T \in \Delta^{k-1}}{\text{minimize}} \quad \frac{1}{m} \sum_{j=1}^m L_{\text{eval},j}(f_T) \\ & \text{s.t. } L_{\text{eval},j}(f_t) = L_{\text{eval},j}(f_{t-1})(1 - A_{:,j}^\top p_{t-1}) \\ & \quad \forall t \in [T], j \in [m] \end{aligned}$$

Skill-It optimization problem

Derive Skill-It's update rule using online mirror descent:

- Regularize with negative entropy (similar to multiplicative weights)
- Proportion (unnormalized) for skill $s_{train, i}$ at time $t+1$:

$$p_{t+1}^i = p_t^i \exp \left(\eta \sum_{j=1}^m A_{ij} L_{\text{eval}, j}(f_t) \right)$$

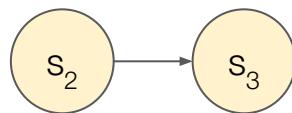
- $\eta > 0$ learning rate
- Intuitively: adjust weight on skill $s_{train, i}$ based on the losses of skills that $s_{train, i}$ influences
 - Key assumption: more data = validation loss goes down

Skill-It optimization problem

Toy example:

- $k = 3$ skills, $S_{\text{train}} = S_{\text{eval}}$
- $\text{Eta} = 0.2$
- Skills graph adjacency matrix:

	S_1	S_2	S_3
S_1	1	0	0
S_2	0	1	0.5
S_3	0	0	1

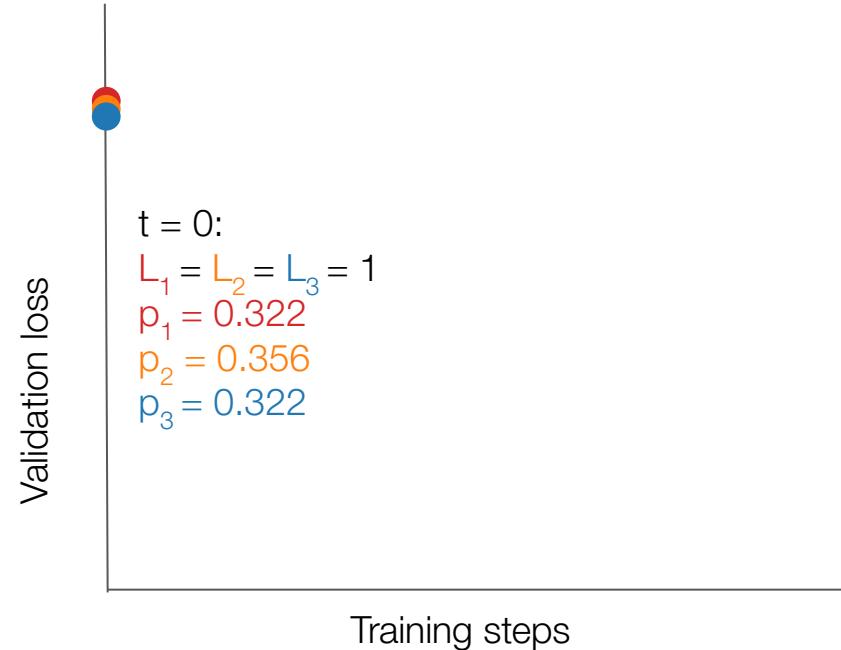
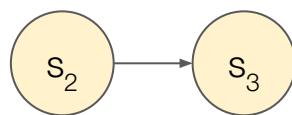


Skill-It optimization problem

Toy example:

- $k = 3$ skills, $S_{\text{train}} = S_{\text{eval}}$
- $\text{Eta} = 0.2$
- Skills graph adjacency matrix:

	s_1	s_2	s_3
s_1	1	0	0
s_2	0	1	0.5
s_3	0	0	1

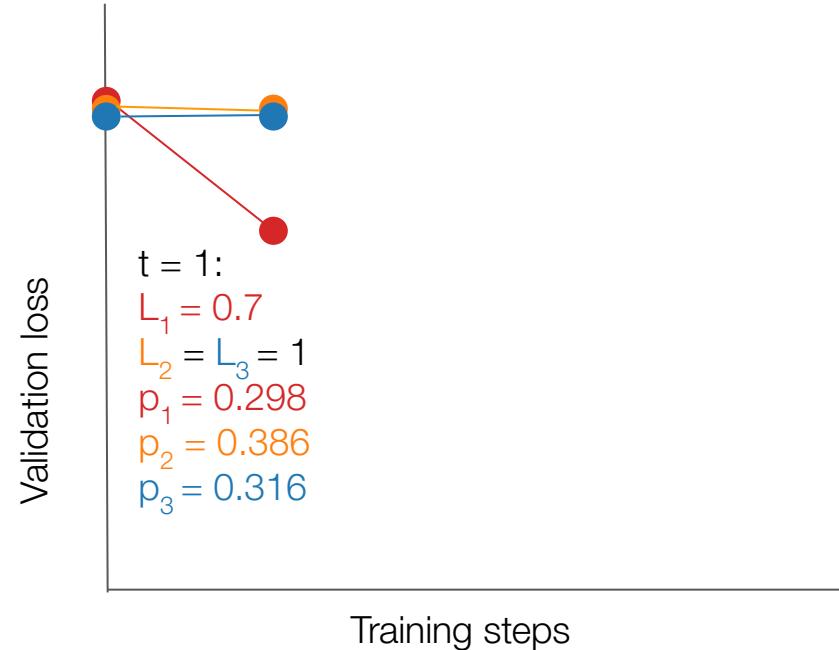
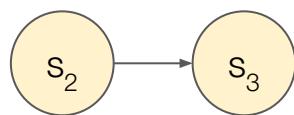


Skill-It optimization problem

Toy example:

- $k = 3$ skills, $S_{\text{train}} = S_{\text{eval}}$
- $\text{Eta} = 0.2$
- Skills graph adjacency matrix:

	S_1	S_2	S_3
S_1	1	0	0
S_2	0	1	0.5
S_3	0	0	1

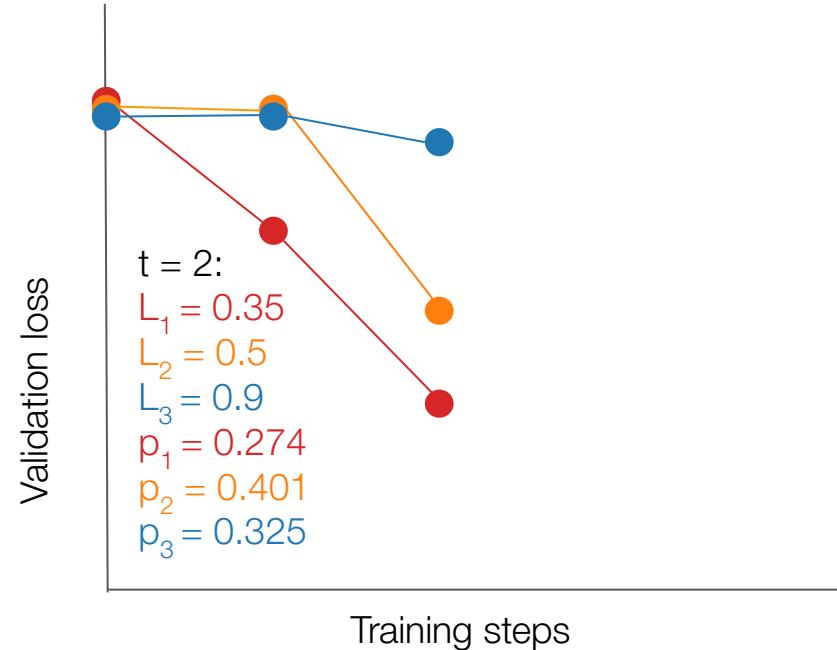
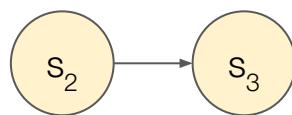


Skill-It optimization problem

Toy example:

- $k = 3$ skills, $S_{\text{train}} = S_{\text{eval}}$
- $\text{Eta} = 0.2$
- Skills graph adjacency matrix:

	S_1	S_2	S_3
S_1	1	0	0
S_2	0	1	0.5
S_3	0	0	1

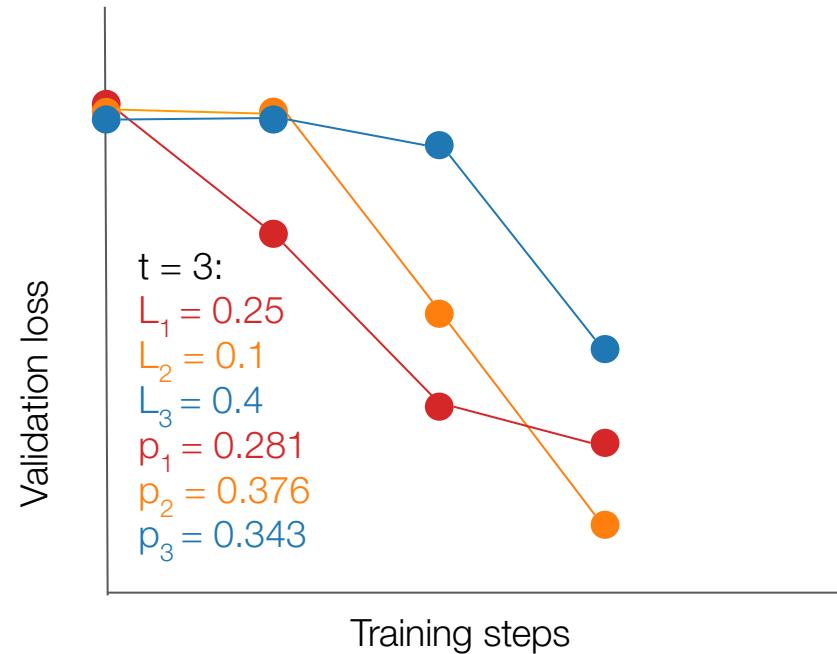
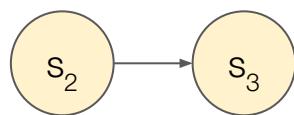


Skill-It optimization problem

Toy example:

- $k = 3$ skills, $S_{\text{train}} = S_{\text{eval}}$
- Eta = 0.2
- Skills graph adjacency matrix:

	s_1	s_2	s_3
s_1	1	0	0
s_2	0	1	0.5
s_3	0	0	1

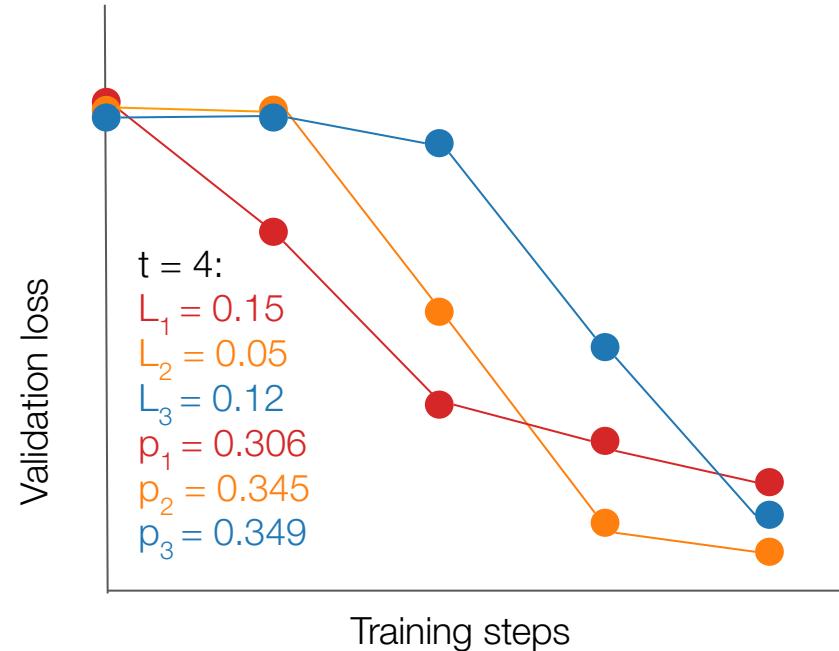
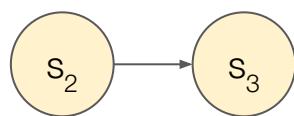


Skill-It optimization problem

Toy example:

- $k = 3$ skills, $S_{\text{train}} = S_{\text{eval}}$
- Eta = 0.2
- Skills graph adjacency matrix:

	s_1	s_2	s_3
s_1	1	0	0
s_2	0	1	0.5
s_3	0	0	1

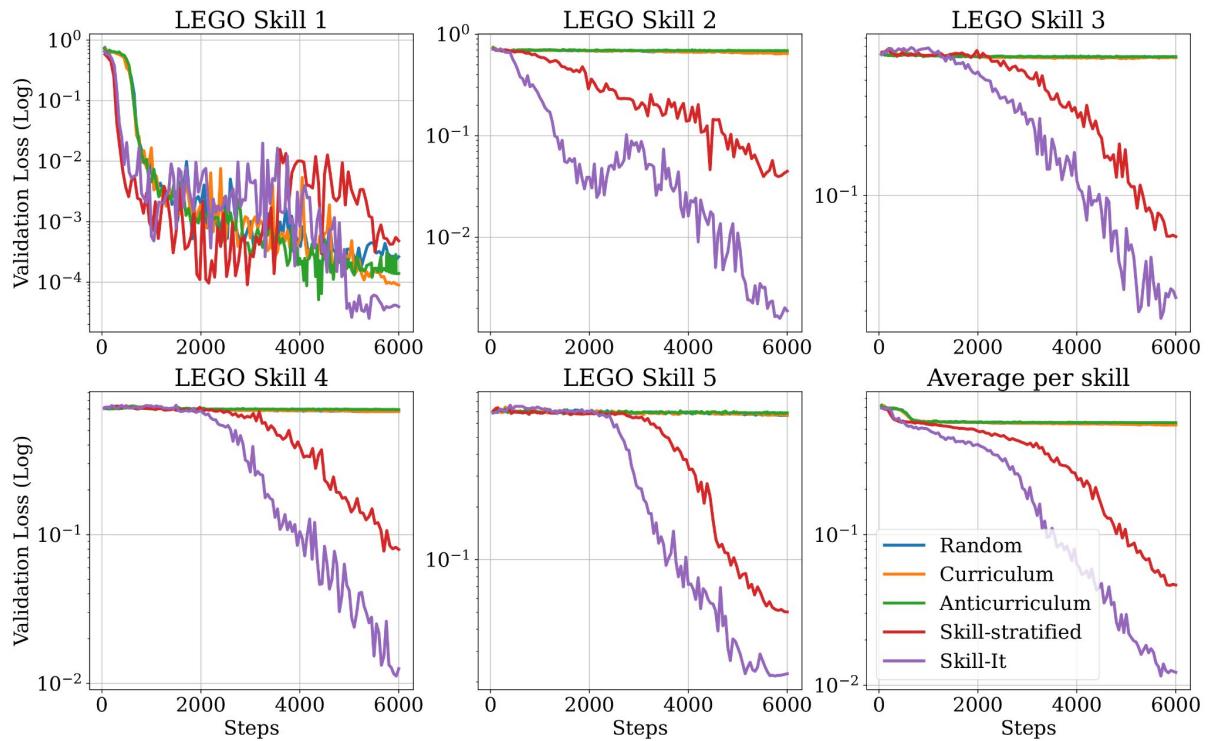
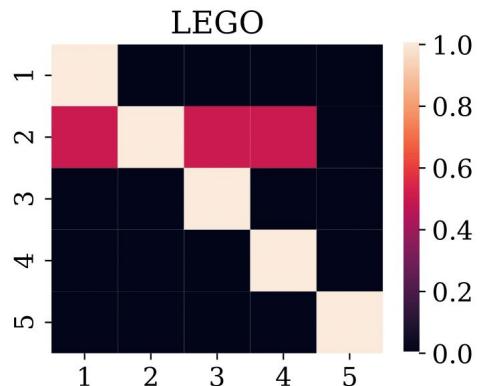


Results

Experiments Overview

- Evaluate Skill-It and skill-stratified sampling in three settings: continual pre-training, fine-tuning, out-of-domain
- Baselines:
 - Random sampling over training corpus or target skill (for fine-tuning)
 - For continual pre-training: curriculum learning at instance level and skill level
- Model: GPT-Neo-125M
 - Experiments where we learn skills graph on 125M and use it on GPT-Neo-1.3B
 - 3B parameter model with together.ai
- Datasets:
 - LEGO
 - Addition
 - Natural Instructions (various subsets)
 - RedPajama (together)

LEGO continual pre-training results



LEGO pre-training results

	Skill 1	Skill 2	Skill 3	Skill 4	Skill 5	Average
Random	100.0 \pm 0.0	54.2 \pm 5.9	58.0 \pm 3.1	48.0 \pm 6.3	54.4 \pm 7.3	62.9 \pm 3.5
Curriculum	100.0 \pm 0.0	60.0 \pm 10.6	55.2 \pm 5.8	51.2 \pm 6.3	51.8 \pm 6.1	63.6 \pm 3.6
Anticurriculum	100.0 \pm 0.0	53.4 \pm 2.3	49.0 \pm 4.8	48.2 \pm 6.4	56.0 \pm 5.7	61.3 \pm 2.2
Skill-stratified	100.0 \pm 0.0	98.2 \pm 1.8	98.2 \pm 1.3	97.8 \pm 1.6	98.2 \pm 1.3	98.5 \pm 0.9
Skill-curriculum	100.0 \pm 0.0	75.2 \pm 30.1	52.2 \pm 3.7	51.0 \pm 4.6	54.4 \pm 3.1	66.6 \pm 7.7
Skill-anticurriculum	100.0 \pm 0.0	90.2 \pm 8.1	88.2 \pm 8.3	73.2 \pm 12.2	62.4 \pm 9.4	82.8 \pm 4.9
SKILL-IT	100.0 \pm 0.0	99.2 \pm 0.8	99.0 \pm 1.0	99.4 \pm 0.5	99.6 \pm 0.5	99.4 \pm 0.2

Table 5: Results on accuracy per skill (binary classification) for LEGO pre-training experiment, averaged over 5 random seeds.

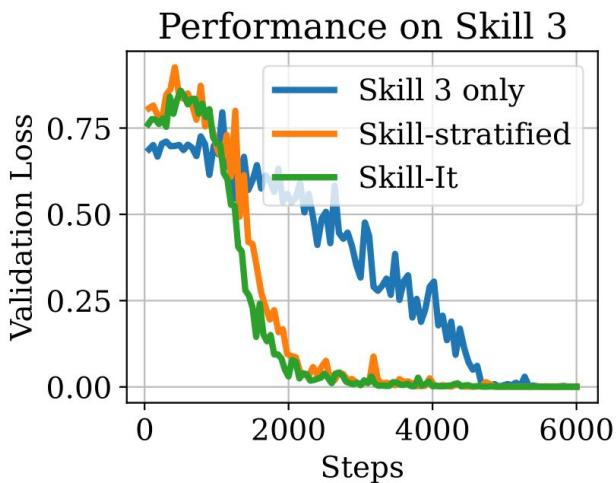
Natural instructions continual pre-training results

- Subset of 23 task categories (> 1M input/output pairs to select from)

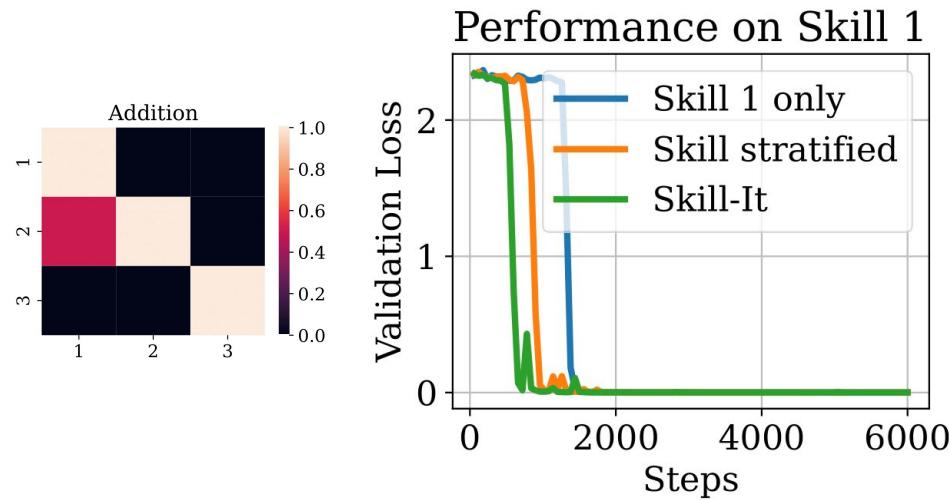
Method	Random	Curriculum	Anticurriculum	Skill curriculum	Skill anticuirlulum	Skill-stratified	Skill-it
Average val loss per skill	2.173 ± 0.028	2.307 ± 0.025	2.366 ± 0.026	2.304 ± 0.031	2.317 ± 0.052	2.115 ± 0.027	2.103 ± 0.032

LEGO and Addition fine-tuning results

LEGO: $S_{\text{eval}} = s_3$

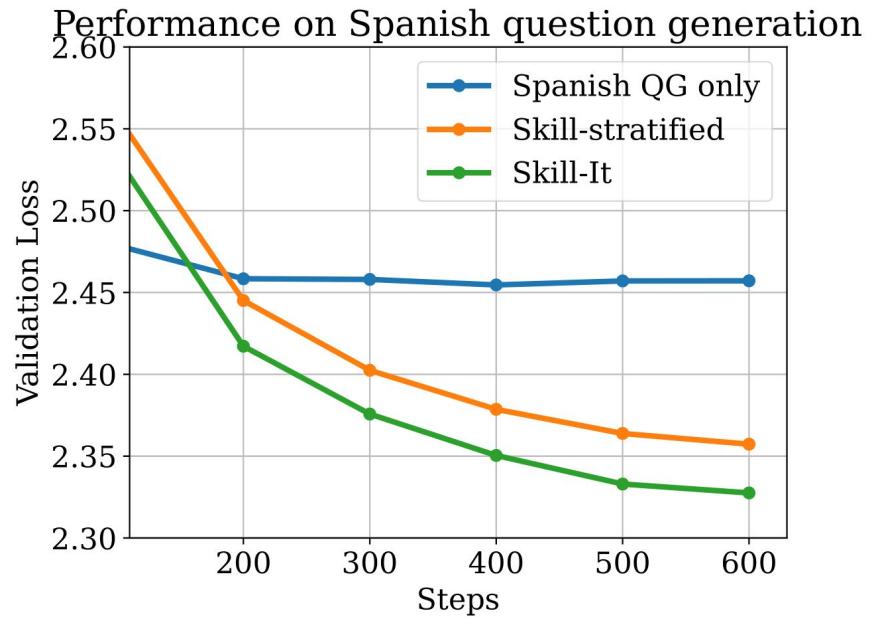
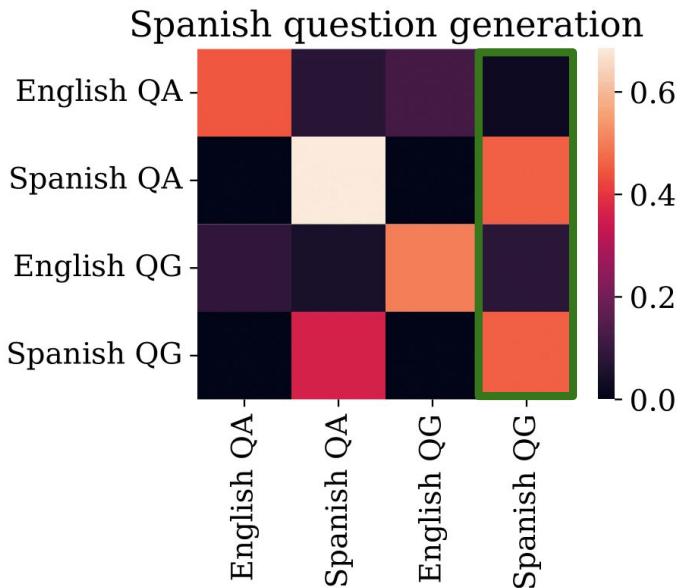


Addition: $S_{\text{eval}} = s_1$



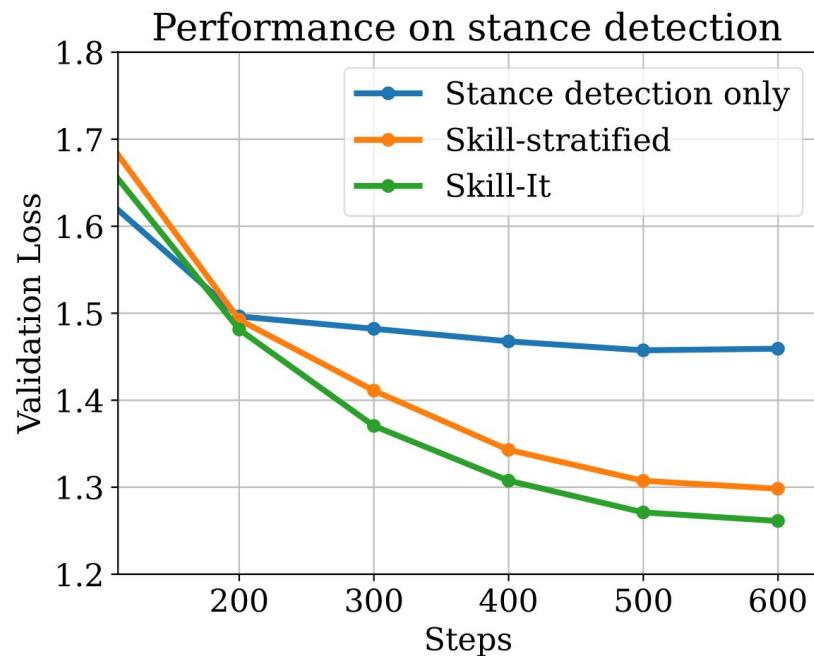
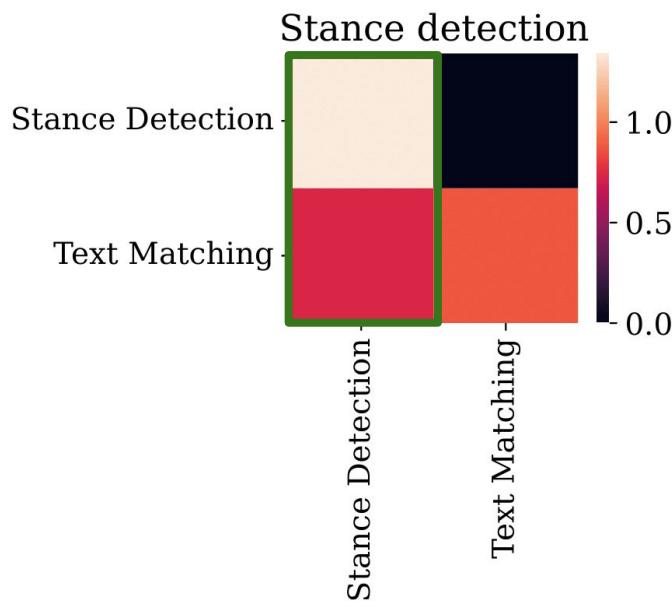
Natural Instructions fine-tuning results

Spanish Question Generation



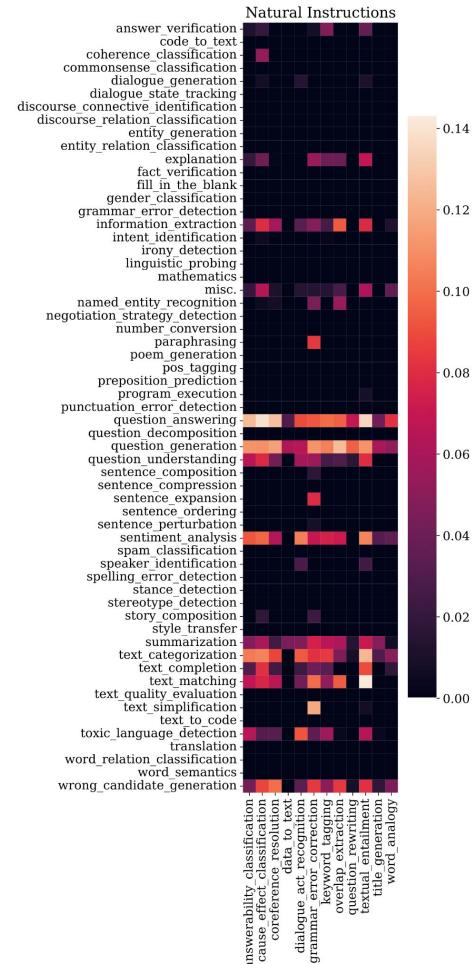
Natural Instructions fine-tuning results

Stance Detection



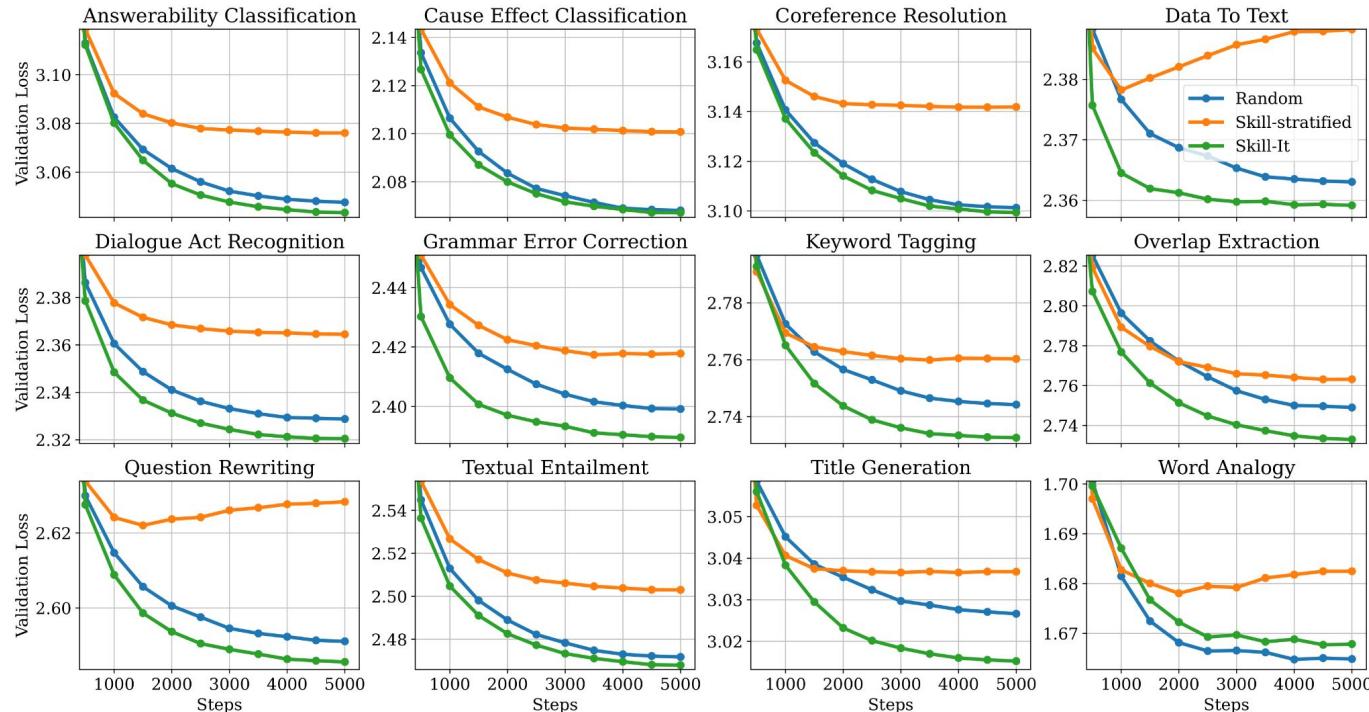
Natural instructions out-of-domain results

- Use 59 task categories from training split as training skill set
- Use 12 task categories from testing split as evaluation skill set



Natural instructions out-of-domain results

Skill-It outperforms baselines on 11/12 task categories



Large-scale case study: RedPajama¹

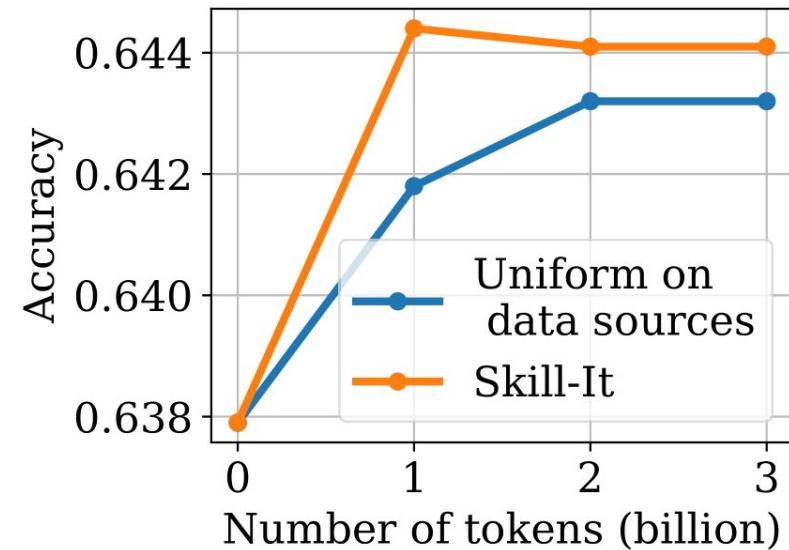
- Training skill set: RedPajama's data sources (ArXiv, Books, C4, CommonCrawl, GitHub, StackExchange, Wikipedia)
- Evaluation skill set: Language Model Evaluation Harness² (ARC Challenge, Bool Q, Copa, HellaSwag, LAMBADA, PIQA, Winogrande)
- Continually pre-train a 3B parameter model already trained on 1T tokens for 3B additional tokens
- Set $T = 1$ (one round/static p_t)

[1] Together, 2023.

[2] Gao et. al. A framework for few-shot language model evaluation, 2021.

Large-scale case study: RedPajama

RedPajama source	SKILL-IT mixture
ArXiv	0.1370
Books	0.0437
C4	0.4195
CommonCrawl	0.0732
GitHub	0.189
StackExchange	0.0892
Wikipedia	0.0484

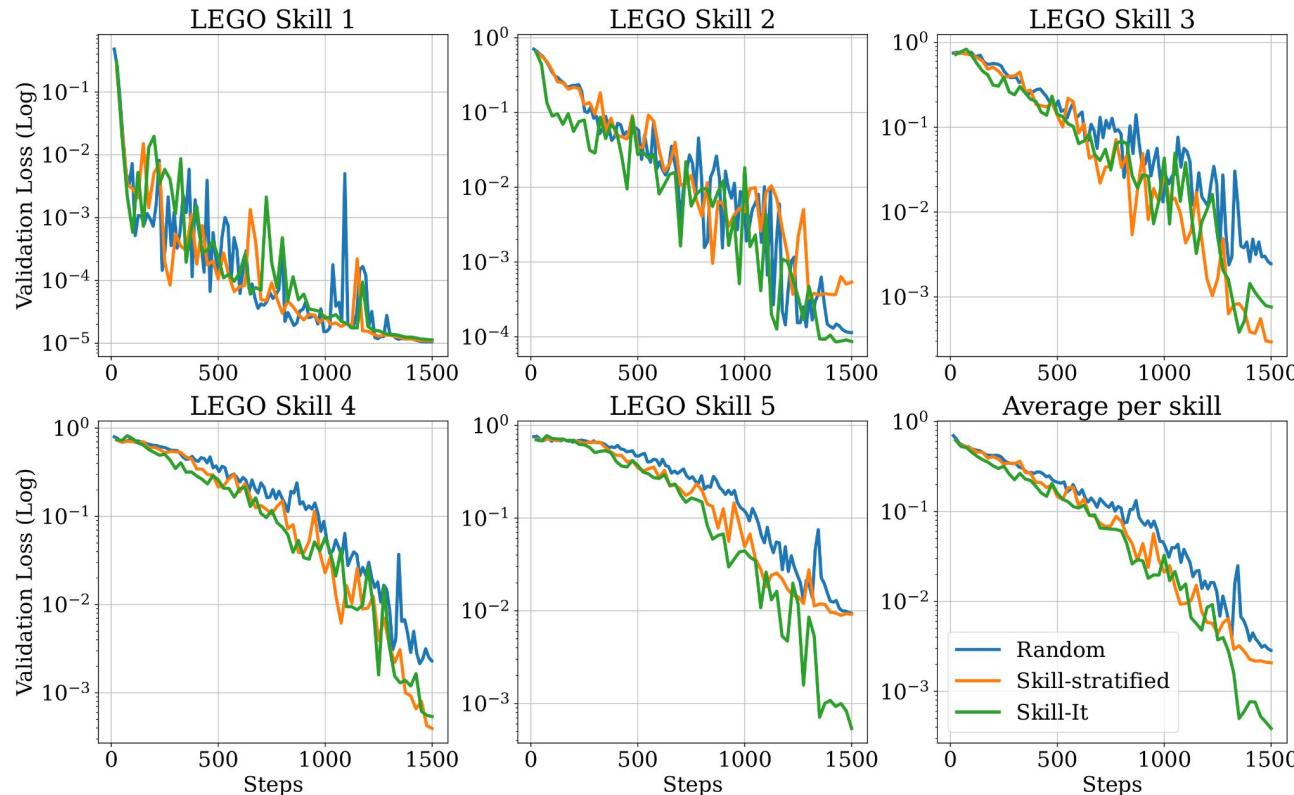


Transfer experiments

Can we learn the skills graph on a smaller model and use it on a larger model?

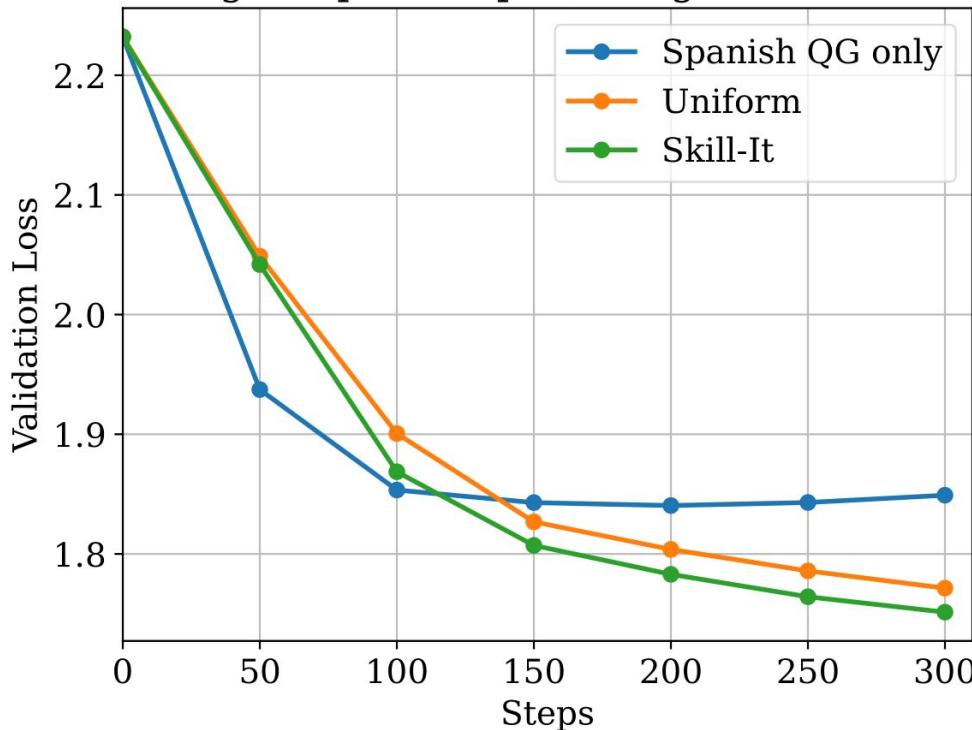
- Smaller model: GPT-Neo-125M
- Larger model: GPT-Neo-1.3B

Transfer experiments: LEGO continual pre-training



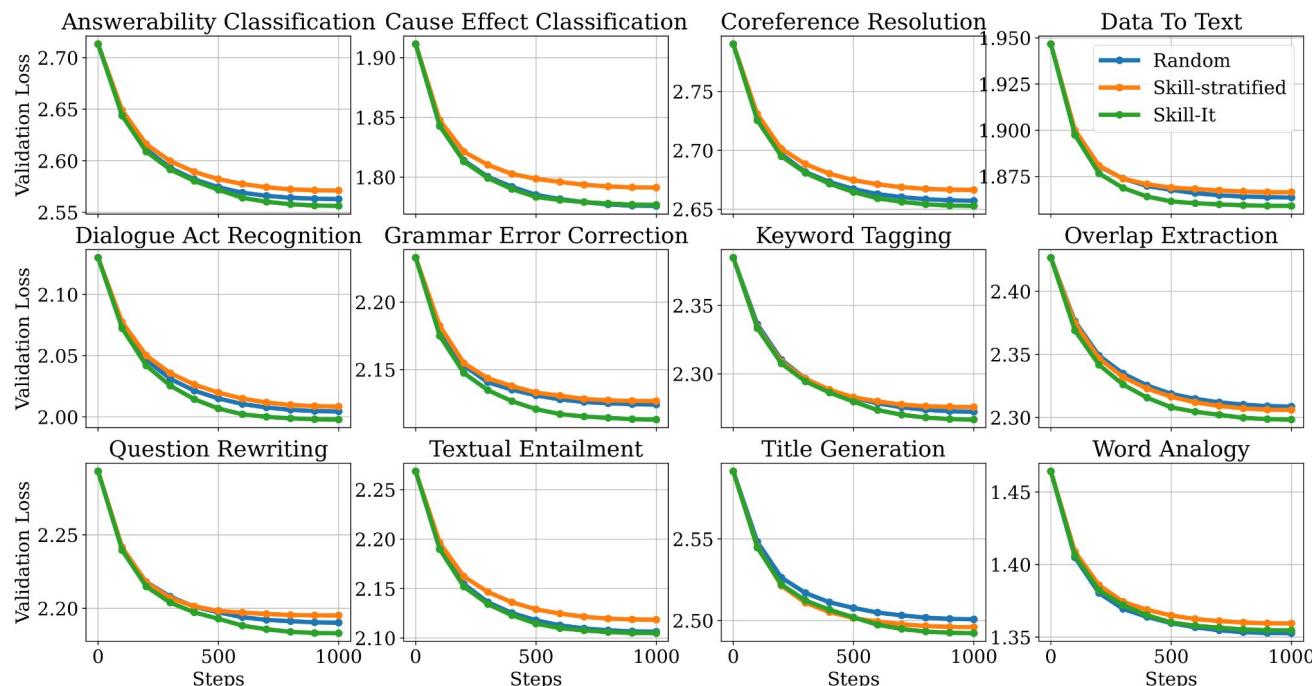
Transfer experiments: Natural Instructions fine-tuning

Targeted learning of Spanish question generation (125M -> 1.3B)



Transfer experiments: Natural Instructions out-of-domain

Skill-It outperforms baselines on 10/12 task categories.



Discussion

Summary

- Given a fixed budget, we want to know what data to train on in order to induce various capabilities in an LM
- We introduce a skills-based framework for understanding how LMs learn and for selecting training data
 - There exist ordered skill sets in the data that the model learns most efficiently in some particular order
 - We propose skill-It, an online algorithm for sampling from skills that uses the skills graph

Discussion

- The definition of skills and skills graph: intuition is not always correct!
- “Are skills just [tasks/data sources/etc.]?”
 - Not always. Skill are a **model-centric** concept
 - Bad set of candidate skills → empty or complete skills graph, Skill-It won’t offer much gain in performance
 - **Discovering ordered skill set from data is hard** (see next slide).

Discussion

- The definition of skills and skills graph: intuition is not always correct!
- “Are skills just [tasks/data sources/etc.]?”
 - Not always. Skill are a **model-centric** concept
 - Bad set of candidate skills → empty or complete skills graph, Skill-It won’t offer much gain in performance
 - **Discovering ordered skill set from data is hard** (see next slide).
- “Can we use something simpler to get the skills graph?”
 - Edges are not always obvious (for example, LEGO)
 - Why does LEGO skill 2 help skill 3 and 4 while skill 3 does not help skill 4?
 - Tried using distances in embedding space/EMD, extremely noisy

Discussion

- The definition of skills and skills graph: intuition is not always correct!
- “Are skills just [tasks/data sources/etc.]?”
 - Not always. Skill are a **model-centric** concept
 - Bad set of candidate skills → empty or complete skills graph, Skill-It won’t offer much gain in performance
 - **Discovering ordered skill set from data is hard** (see next slide).
- “Can we use something simpler to get the skills graph?”
 - Edges are not always obvious (for example, LEGO)
 - Why does LEGO skill 2 help skill 3 and 4 while skill 3 does not help skill 4?
 - Tried using distances in embedding space/EMD, extremely noisy
- How to improve the Skill-It algorithm
 - Graph is *dynamic* and needs to be updated - how much can we do on the fly?

Future work: skills discovery

- Can we discover skills by clustering?
- Task: recover LEGO skills by clustering:
 - Embeddings
 - Gradients - in progress
 - Validation losses across multiple training runs
 - Points that have the same shifts/behaviors in validation loss as a function of training data tend to belong to the same skill

Cluster method	Accuracy
Pretrained embedding of last token	24.8 ± 0.5
Pretrained embedding of average token	25.2 ± 1.1
Trained model embedding of last token	38.4 ± 0.8
Sentence-BERT embedding	23.9 ± 0.7
Losses over multiple runs	61.0 ± 1.6

Thank you! Questions?

Contact: mfchen@stanford.edu

ArXiv: <https://arxiv.org/abs/2307.14430>

